

欢迎大家参加
IBM开源技术微讲堂
Kubernetes系列之 —— 初探

IBM开放技术研究院
2017.10.19



IBM开放技术研究院

- 专注于开源技术与开放标准的开发
 - 公众号：ibmopentech
- 微讲堂系列活动
 - 每周四晚8点
 - WebEx和斗鱼同步直播
 - <http://ibm.biz/opentech-ma>
 - 已经完成了
 - Open Stack系列
 - Cloud Foundry系列
 - Container技术系列
 - Hyperledger系列
 - Apache Open Whisk系列



Kubernetes系列

- 10月19日 Kubernetes初探
- 10月26日 上手Kubernetes：基本概念、安装和命令行工具kubectl
- 11月2日 Kubernetes的资源调度
- 11月9日 Kubernetes的运行时：kubelet
- 11月16日 Kubernetes的网络管理
- 11月23日 Kubernetes的存储管理
- 11月30日 Kubernetes的日志与监控
- 12月7日 Kubernetes的应用部署
- 12月14日 扩展Kubernetes生态：Service Catalog的概念与应用
- 12月21日 Kubernetes的企业实践



第一讲：Kubernetes初探



- Doug Davis

就职于IBM Digital Business Group的开源与开放标准项目组，专注于开源云计算技术（例如Docker与Kubernetes）的开发与研究。在他的职业生涯中，曾参与过许多开源项目与开放标准的研发工作，包括Cloud Foundry、Apache Axis、CIM以及大部分SOAP相关的标准等。目前，他还代表IBM在许多国际标准组织中，如DMTF、OASIS、W3C等，担任要职

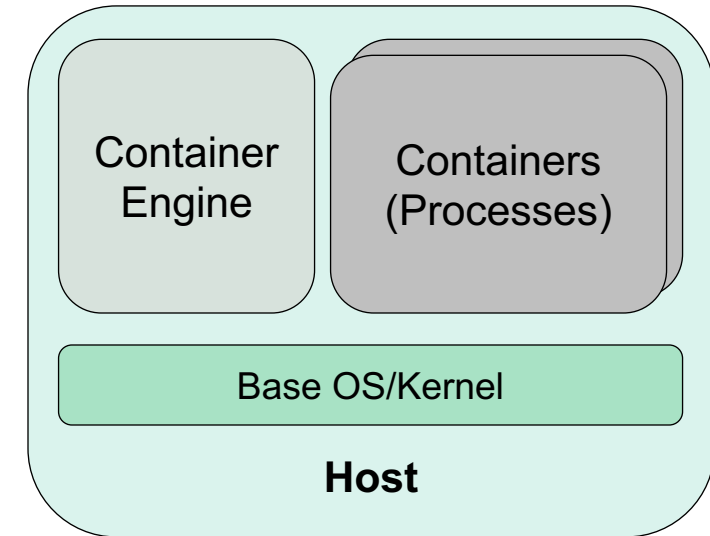


Today's Agenda

- What is Kubernetes?
- How was Kubernetes created?
- Where is the Kubernetes community?
- Technical overview
- What's the current status of Kubernetes?
- What is next for the Kubernetes community?
- Kubernetes at IBM
- Let's Get Started
- Call to Action

What are Containers?

- Container: a set of processes run in isolation
- Each container gets its own:
 - PID, User, UTS, Mount Points, Network Stack, etc...
 - And its own view of the filesystem
- Very similar to VMs
 - But process based – run just the app itself, nothing else
 - No operating system, just the Linux kernel files are available
- Benefits:
 - Smaller footprint – just the application's files
 - **Faster start-up times – just starting the exe – not even the OS**
 - **Milliseconds vs minutes**
 - All adds up to better resource utilization at faster scale



What is Kubernetes?

Enterprise Level Container Orchestration



- Provision, manage, scale applications (containers) across a cluster
- Manage infrastructure resources needed by applications
 - Volumes
 - Networks
 - Secrets
 - And many many many more...
- Declarative model
 - Provide the "desired state" and Kubernetes will make it happen
- What's in a name?
 - Kubernetes (K8s/Kube): "Helmsman" in ancient Greek

How was Kubernetes created?

Google Bringing its Cloud-Scale Expertise to OSS

- Based on Google's internal "Borg" project
- Not just Open Source, but Open Governance!
 - Working very hard to ensure that it is not a "Google" project
 - For example, it was the reason the CNCF was created
 - Pushing for non-Googlers to be in key leadership roles
- Quickly attracted the attention & support of others
 - Looking for alternatives to Docker
 - E.g. RedHat, CoreOS, Deis/EngineYard (now Microsoft), IBM
- One of the fastest growing OSS projects
 - Which has brought many scaling challenges to deal with
 - E.g. had to have "stability" releases to deal with the rapid pace of changes



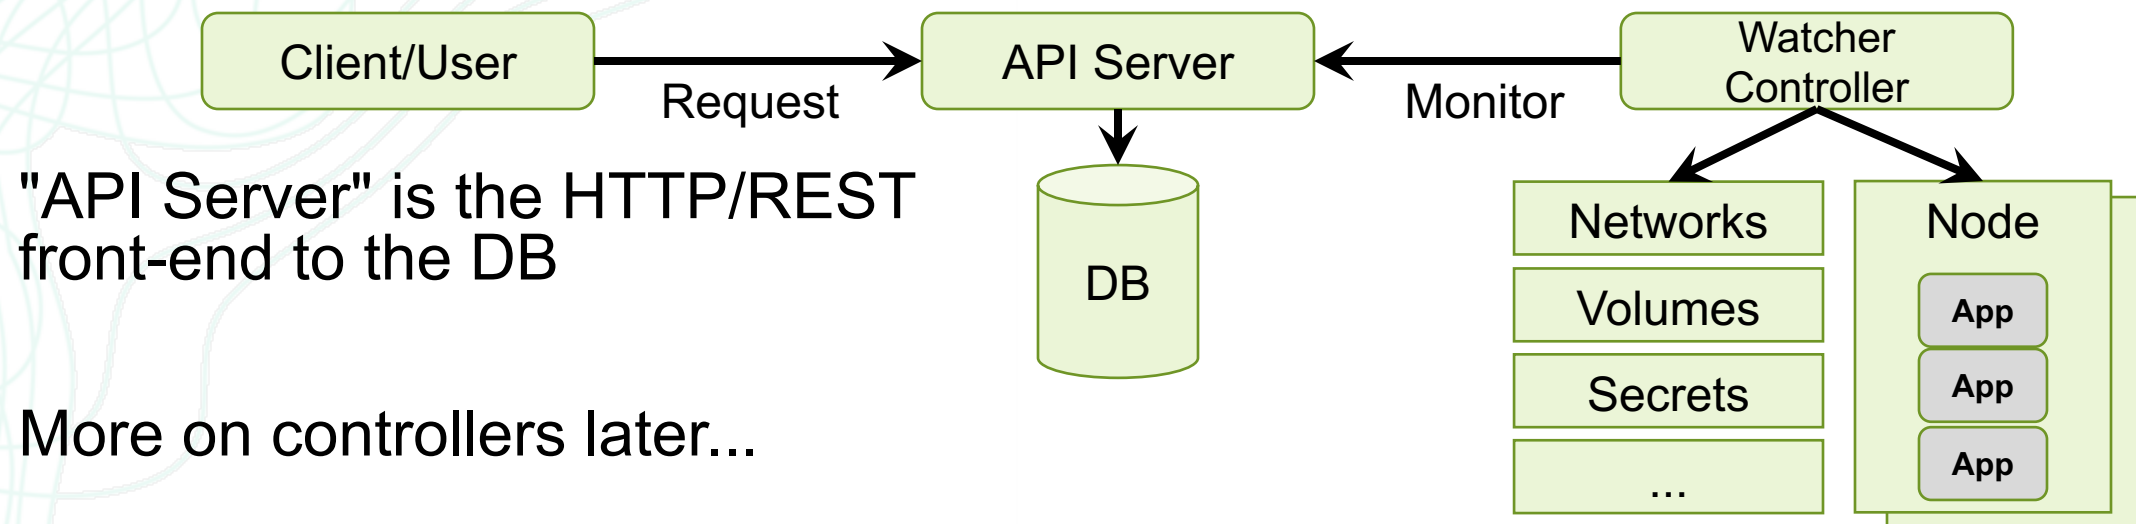
Where is the Kubernetes community?

- Main source entry point: <https://github.com/kubernetes/>
- Communications
 - <https://github.com/kubernetes/community/blob/master/communication.md>
- Mailing list / google groups:
 - Devs: kubernetes-dev@googlegroups.com
 - Users: kubernetes-users@googlegroups.com
 - Weekly Community Meeting
 - Conferences (CNCF-con & KubeCon)
- Lots of SIGs (special interest groups) for focused areas/functionality
 - <https://github.com/kubernetes/community/blob/master/sig-list.m>
 - Each with their own slack channel, mailing lists, regular calls, ...
- Leaders: Google, RedHat, CoreOS, Microsoft (Deis), IBM, Huawei



Kubernetes: Technical Overview

- At its core, Kubernetes is a database (etcd).
With "watchers" & "controllers" that react to changes in the DB.
The controllers are what make it Kubernetes.
This pluggability and extensibility is part of its "secret sauce".
- DB represents the user's desired state.
Watchers attempt to make reality match the desired state



Kubernetes: Resource Model

A resource for just about any purpose

- Config Maps
- Daemon Sets
- Deployments
- Events
- Endpoints
- Ingress
- Jobs
- Nodes
- Namespaces
- **Pods**
- Persistent Volumes
- Replica Sets
- Secrets
- Service Accounts
- Services
- Stateful Sets, and more...

- Kubernetes aims to have the building blocks on which you build a cloud native platform.
- Therefore, the internal resource model **is** the same as the end user resource model.

Key Resources

- Pod: set of co-located containers
 - Smallest unit of "code" deployment
- Application: undefined, but is a set of pods
 - Several types of resources to help manage them
 - Replica Sets, Deployments, Stateful Sets, ...
- Services & Endpoints
 - Define how to expose your app
 - Query based selector to choose which pods apply

Kubernetes: Technical Overview

- The user directly manipulates resources via json/yaml

```
$ kubectl (create|get|apply|delete) -f myResource.yaml
```

- Some attempts to soften the UX:

```
$ kubectl scale ...
```

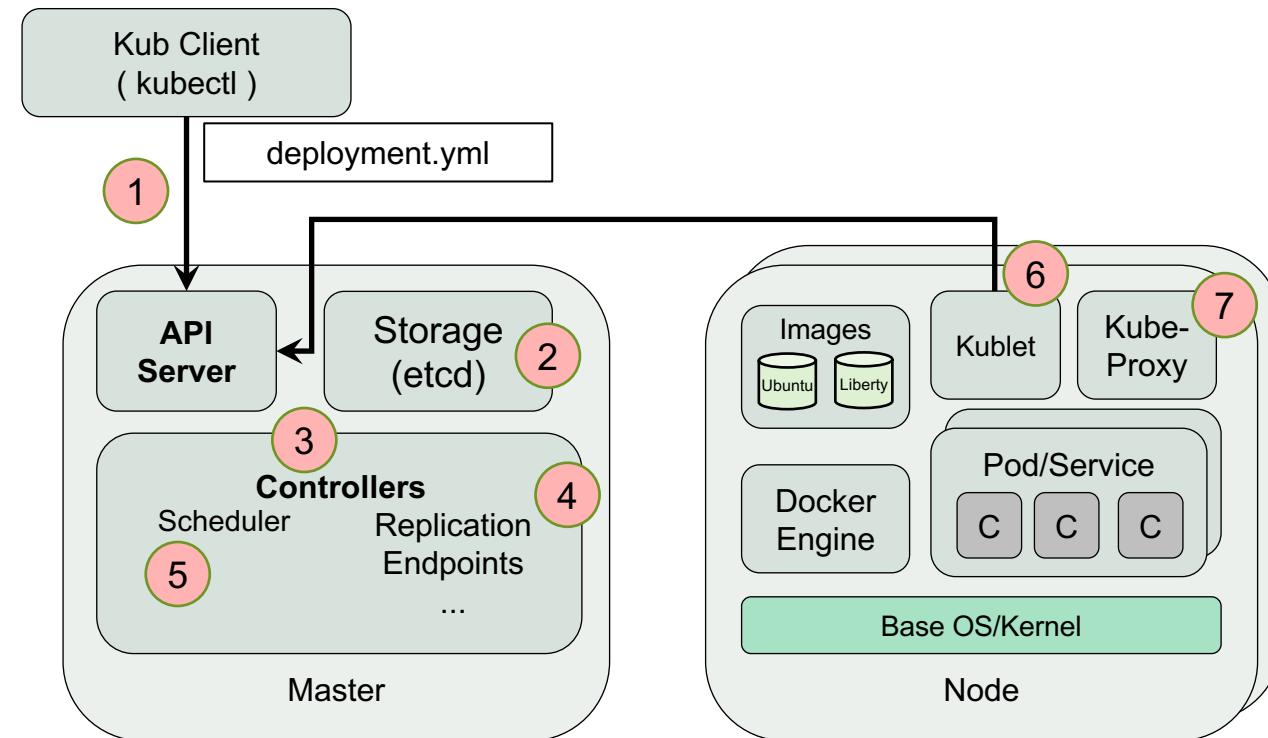
```
$ kubectl run ...
```

```
$ kubectl annotate ...
```

- But those are limited and not the norm

Kubernetes: Putting it all together...

1. User via "kubectl" deploys a new application
2. API server receives the request and stores it in the DB (etcd)
3. Watchers/controllers detect the resource changes and act upon it
4. ReplicaSet watcher/controller detects the new app and creates new pods to match the desired # of instances
5. Scheduler assigns new pods to a kubelet
6. Kubelet detects pods and deploys them via the container runing (e.g. Docker)
7. Kube-proxy manages network traffic for the pods – including service discovery and load-balancing



What is the current status of Kubernetes?

- V1.7 - celebrated its two year birthday (July 21, 2017) since v1.0
- Some newer features
 - Custom Resource Definitions replaces Third Party Resources
 - Network Policy API is "stable" – control for pod->pod communication
 - Encryption for data at rest in Secrets
 - "Local Storage" persistent volume type was added (alpha)
 - API Aggregation
 - External admission controllers
 - Role Based Access Control (RBAC) is in beta
 - Kubeadm (tool for deploying kube) is in beta
 - Node affinity/anti-affinity, taints, tolerations for scheduling

What's next for the Kubernetes community?

Explosive Growth – a good problem to have

- Kubernetes is going through some growing pains
 - Rapid growth (code & community)
 - Finds itself needing to "slow down" at times to ensure stability
 - Stability releases
 - Lock down amount and rate of changes
- Aside from being a CN platform, K8s is a promoter of interop
 - Container Storage Interface (CSI) w/ Docker, CF
 - Container Networking Interface (CNI) w/ Docker, CF
 - Open Service Broker API (SIG-ServiceCatalog) w/ CF
 - Istio w/ CF

Kubernetes at IBM

Preferred Container Orchestration

- Offerings / Plans
 - Bluemix Kubernetes Service – Docker containers orchestrated by K8s
 - ICp – IBM Cloud Private
 - Watson is leveraging Kubernetes to hosting its infrastructure
- Key Development Activities
 - Service Catalog (co-lead)
 - Contributor Experience
 - Networking & Istio (co-lead)
 - ContainerD integration (co-lead)
 - Storage
 - Performance

Kubernetes: Let's Get Started

- Development Guide
 - Getting Started Guide from IBM's Mike Brown
 - Kubernetes Source Code Tour from IBM's Brad Topol
 - Help: Ask on the appropriate SIG Slack channel
- Looking for work:
 - Backlog of issues – many many open issues
 - Contributor experience, testing, "process" related activities
 - Find us on the Slack channel and ask questions!
- Journeys
 - <https://developer.ibm.com/code/journey/run-gitlab-kubernetes/>
 - <https://developer.ibm.com/code/journey/deploy-microprofile-java-microservices-on-kubernetes/>
 - <https://developer.ibm.com/code/events/manage-microservices-traffic-using-istio/>

Future Kubernetes Classes

- Start K8s Journey - K8s concepts and kubectl
- Connect to the world - K8s service catalog
- What makes Kubernetes smart? Scheduling in Kubernetes
- Everything is well connected and located - Kubernetes network
- Stateless vs Stateful? Kubernetes storage
- Keep healthy - Kubernetes logging and monitoring
- Play with applications - Kubernetes Helm and Charts
- What we can do for you - IBM ICp

Q & A

Call to Action

- [Get involved!](#)
- [Join a meetup](#)
- [Attend a conference](#)
- [Join the discussion](#)
- [Contribute code](#)

Backup

Watchers and Controllers

Some examples:

1. Replica Set Controller

- Verifies the correct number of "pod" instances are active
 - ReplicaSet is a scalable set of pods

2. Scheduler

- Watches for new pods and assigns them to a "kubelet"
 - Pod is a group of containers that share lifecycle and container resources

3. Kubelet

- Watches for pods to be assigned to it, then deploys the pod
 - Kubelet manages the pods/containers running on a host

